


PYTORCH



Ben

INTRODUCTION

- Pytorch is a deep learning framework that puts Python first.
- Based on Torch but write in Python.
- Tensors and Dynamic neural networks in Python with strong GPU acceleration.

facebook



 NVIDIA



ParisTech
INSTITUT DES SCIENCES ET TECHNOLOGIE
PARIS INSTITUTE OF TECHNOLOGY

Carnegie
Mellon
University

UNIVERSITÉ
PIERRE & MARIE CURIE
SCIENCE À PARIS

 Digital
Reasoning

 Stanford
University



Inria



COMPONENTS

<code>torch</code>	a Tensor library like NumPy, with strong GPU support
<code>torch.autograd</code>	a tape-based automatic differentiation library that supports all differentiable Tensor operations in torch
<code>torch.nn</code>	a neural networks library deeply integrated with autograd designed for maximum flexibility
<code>torch.multiprocessing</code>	Python multiprocessing, but with magical memory sharing of torch Tensors across processes. Useful for data loading and Hogwild training.
<code>torch.utils</code>	DataLoader, Trainer and other utility functions for convenience
<code>torch.legacy(.nn/.optim)</code>	legacy code that has been ported over from torch for backward compatibility reasons



FEATURES

- A GPU-ready Tensor library
- Dynamic Neural Networks: Tape based autograd
- Python first
- Fast
- Extensions without pain



GPU ACCELERATION

- Based on NVIDIA CUDA
- Use CUDA semantics then run on GPU
- Support multi-GPU on a single machine.

```
dtype = torch.FloatTensor  
# dtype = torch.cuda.FloatTensor # Uncomment this to run on GPU
```



DYNAMIC NEURAL NETWORKS

Most frameworks such as TensorFlow, Theano, Caffe and CNTK have a static view of the world. [A graph is created on the fly](#)

- One has to build a neural network, and reuse the same structure again and again.
- Changing the way the network behaves means that one has to start from scratch.

With PyTorch, we use a technique called Reverse-mode auto-differentiation, which allows you to change the way your network behaves arbitrarily with zero lag or overhead.

```
from torch.autograd import Variable

x = Variable(torch.randn(1, 10))
prev_h = Variable(torch.randn(1, 20))
W_h = Variable(torch.randn(20, 20))
W_x = Variable(torch.randn(20, 10))
```



COMPARISON

- TensorFlow is a safe bet for most projects. Not perfect but has huge community, wide usage.
- I(Fei-Fei Li) think Pytorch is **best** for research. However still new, there can be rough patches.
- Use TensorFlow for one graph over many machines
- Consider Caffe, Caffe2, or TensorFlow for production deployment
- Consider TensorFlow or Caffe2 for mobile



CODE COMPARISON

numpy

```
N, D_in, H, D_out = 64, 1000, 100, 10
x = np.random.randn(N, D_in)
y = np.random.randn(N, D_out)
w1 = np.random.randn(D_in, H)
w2 = np.random.randn(H, D_out)
learning_rate = 1e-6
for t in range(500):
    h = x.dot(w1)
    h_relu = np.maximum(h, 0)
    y_pred = h_relu.dot(w2)
    loss = np.square(y_pred - y).sum()
    grad_y_pred = 2.0 * (y_pred - y)
    grad_w2 = h_relu.T.dot(grad_y_pred)
    grad_h_relu = grad_y_pred.dot(w2.T)
    grad_h = grad_h_relu.copy()
    grad_h[h < 0] = 0
    grad_w1 = x.T.dot(grad_h)
    w1 -= learning_rate * grad_w1
    w2 -= learning_rate * grad_w2
```

tensorflow

```
N, D_in, H, D_out = 64, 1000, 100, 10
x = tf.placeholder(tf.float32, shape=(None, D_in))
y = tf.placeholder(tf.float32, shape=(None, D_out))
w1 = tf.Variable(tf.random_normal((D_in, H)))
w2 = tf.Variable(tf.random_normal((H, D_out)))
h = tf.matmul(x, w1)
h_relu = tf.maximum(h, tf.zeros(1))
y_pred = tf.matmul(h_relu, w2)
loss = tf.reduce_sum((y - y_pred) ** 2.0)
grad_w1, grad_w2 = tf.gradients(loss, [w1, w2])
learning_rate = 1e-6
new_w1 = w1.assign(w1 - learning_rate * grad_w1)
new_w2 = w2.assign(w2 - learning_rate * grad_w2)
with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())
    x_value = np.random.randn(N, D_in)
    y_value = np.random.randn(N, D_out)
    for _ in range(500):
        loss_value, _, _ = sess.run([loss, new_w1, new_w2],
                                     feed_dict={x: x_value
```

pytorch

```
dtype = torch.FloatTensor
N, D_in, H, D_out = 64, 1000, 100, 10
x = Variable(torch.randn(N, D_in).type(dtype),
              requires_grad=False)
y = Variable(torch.randn(N, D_out).type(dtype),
              requires_grad=False)
w1 = Variable(torch.randn(D_in, H).type(dtype),
              requires_grad=True)
w2 = Variable(torch.randn(H, D_out).type(dtype),
              requires_grad=True)
learning_rate = 1e-6
for t in range(500):
    y_pred = x.mm(w1).clamp(min=0).mm(w2)
    loss = (y_pred - y).pow(2).sum()
    w1.grad.data.zero_()
    w2.grad.data.zero_()
    loss.backward()
    w1.data -= learning_rate * w1.grad.data
    w2.data -= learning_rate * w2.grad.data
```



DEMO

Classifying Names with a Character-Level RNN

- Preparing the Data
- Turning Names into Tensors
- Creating the Network
- Training the Network
- Evaluating the Results

Arabic.txt

Chinese.txt

Czech.txt

Dutch.txt

English.txt

French.txt

German.txt

Greek.txt

Irish.txt

Italian.txt

Japanese.txt

Korean.txt

Polish.txt

Portuguese.txt

Russian.txt

Scottish.txt

Spanish.txt

Vietnamese.txt

```
print(letter_to_tensor('J'))
```

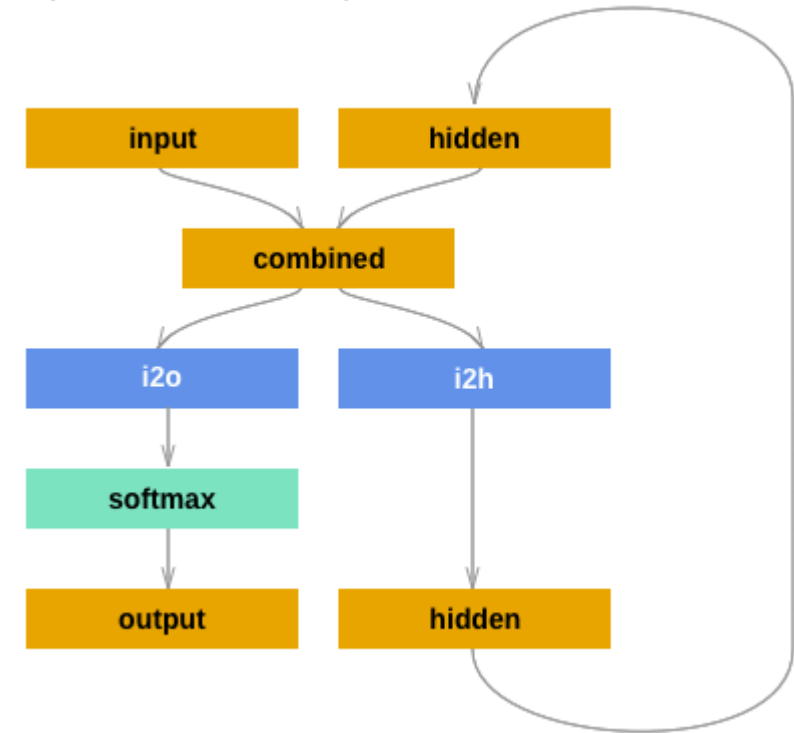
Columns 0 to 12
0 0 0 0 0 0 0 0 0 0 0 0 0

Columns 13 to 25
0 0 0 0 0 0 0 0 0 0 0 0 0


Columns 26 to 38
0 0 0 0 0 0 0 0 0 1 0 0 0

Columns 39 to 51
0 0 0 0 0 0 0 0 0 0 0 0 0

Columns 52 to 56
0 0 0 0 0
[torch.FloatTensor of size 1x57]



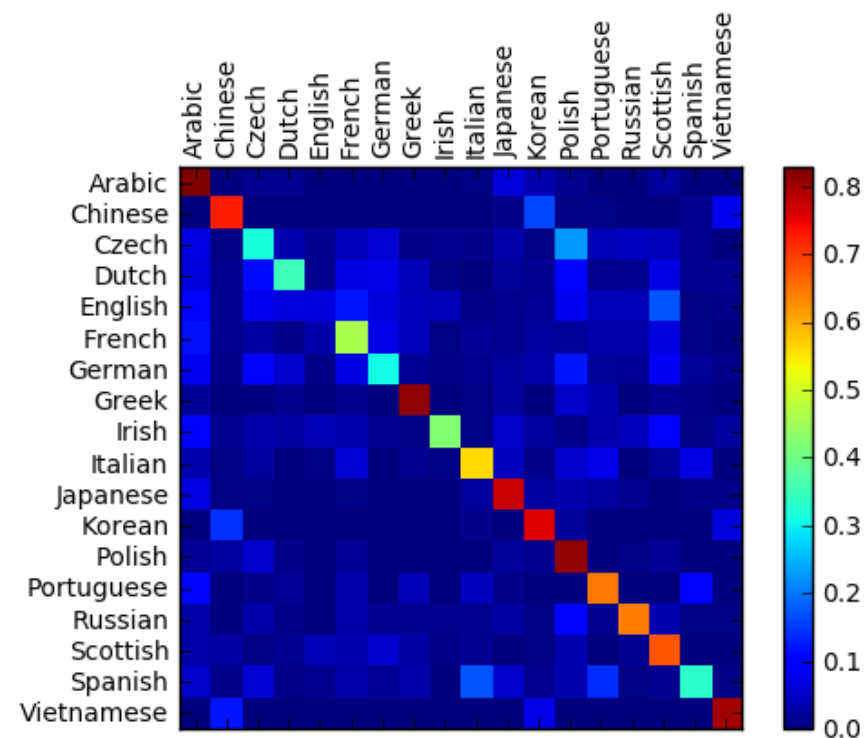
RESULT



> Dovesky
 (-0.87) Czech
 (-0.88) Russian
 (-2.44) Polish

> Jackson
 (-0.74) Scottish
 (-2.03) English
 (-2.21) Polish

> Satoshi
 (-0.77) Arabic
 (-1.35) Japanese
 (-1.81) Polish





HOW TO STUDY IT

- [tutorials](#)
- [pytorch/examples](#)
- [PyTorch doc](#)
- [PyTorch Forums](#)




Thanks !





WORKFLOW (APPLICATION)



➤ 1. Observed data

➤ 2. Find feature

➤ 3. Design algorithms

➤ 4. Verification

➤ 5. Data cleaning

➤ 6. Engineering

➤ 7. Product feedback

➤ 8. To the begin



WORKFLOW (CODING)

A typical training procedure for a neural network is as follows:

- Define the neural network that has some learnable parameters (or weights)
- Iterate over a dataset of inputs
- Process input through the network
- Compute the loss (how far is the output from being correct)
- Propagate gradients back into the network's parameters
- Update the weights of the network, typically using a simple update rule:
$$\text{weight} = \text{weight} + \text{learning_rate} * \text{gradient}$$

